



Information valid as of May 2024. Please reference the AWS Partner: Generative AI on AWS Essentials (Business) Skill Builder course for updated information, generative AI basics, and qualifying questions.

Joint value proposition	
Why build generative AI with AWS	
Customer challenge	AWS solution
Broadest choice of foundation models Customers may need access to a variety of models, and model sizes. They need the flexibility to optimize for cost, performance, and latency for different types of use cases and applications.	AWS was the first to recognize the importance of offering a broad choice of leading FMs from companies including Amazon, AI21 Labs, Anthropic, Cohere, Meta, Mistral AI, and Stability AI via a single API.
Easy to use tools for every type of customer Customers may want to build applications with existing FMs.	Amazon Bedrock allows the customer to select the best FM for their use case and customize applications with their proprietary data in a secure and private environment. We've also built applications with generative AI built inside to help customers without ML expertise quickly take advantage of generative AI and increase employee productivity.
Purpose built infrastructure to help you scale Whether customers are training their own models, customizing models, or running ML applications, they need to train and run inference at scale with an infrastructure that's purpose-built for ML.	AWS invested heavily in purpose-built accelerators, AWS Trainium and AWS Inferentia, to provide the most performant and low-cost infrastructure for generative AI.
A comprehensive set of data services Many customers may have data locked in departmental silos or legacy applications. Their privacy controls may be ad-hoc, and data can be inaccurate and messy.	AWS provides an extensive set of services for managing data. This includes specific-purpose databases, data movement, metadata management, and data governance. With the most comprehensive set of data and AI services, a customer can use its organization's data to securely customize an FM on AWS to build differentiated applications that know your business, data, and customers.
Buyer Persona	Why these persona & conversation starters
Business decision makers (Business executives; VP, head, or director of marketing; finance; operations; customer	<u>Business decision makers</u> This persona has the authority to find and fund technology solutions that can help achieve their business outcomes. They are motivated by improving company performance, creating competitive advantage with innovative technology, and potential AI/ML business results.



<p>service; supply chain; logistics; and legal)</p>	<p>Discovery questions for C-suite business decision makers:</p> <ul style="list-style-type: none">• How would you rank the following in order of importance to choose the right foundation model to build your generative AI application?<ul style="list-style-type: none">○ Performance○ Latency○ Cost○ Accuracy• What customer experiences or other use cases do you expect to transform with generative AI?• How might you improve customer retention by creating more personalized user experiences for your customers?• What is your vision for AI/ML and your company? What are your business goals for AI for this year?• Have you achieved the business outcomes that you were aiming towards with previous AI and machine learning projects?• What success or challenges have you had in working with a technology partner to plan and build your AI/ML capabilities?• What kinds of business operations would you automate first by using AI technology that is capable of generating novel content? <p><u>C-suite CMO and Marketing</u></p> <p>When talking with this persona, you might focus on the following topics: accelerating content production, increasing reach with personalized ads, and quickly creating localized content. To learn more about C-suite marketing discovery questions, choose each of the following flashcards.</p> <p>Discovery questions for C-suite CMO and Marketing:</p> <ul style="list-style-type: none">• How often are you waiting for new media assets to be created to roll out new marketing campaigns?• How would you expand your brand's reach if you had the ability to deliver highly personalized ads?• How would your ability to enter new markets be accelerated if you had the ability to automatically generate marketing material? How would your ability to enter new markets be accelerated if you had the ability to automatically generate marketing material?
<p>Technical decision makers</p> <p>(Chief cloud architect, head or director of data science, analytics, AI and ML, IT executives)</p>	<p><u>Technical decision makers</u> usually seek a deeper understanding of the technology. They want to see product details, use cases, and clear explanations on what AWS can do for generative AI in relation to their individual goals. Their top priorities are increasing innovation, delivering IT projects quickly, aligning IT performance metrics to business outcomes, and cutting overall IT costs.</p> <p>Discovery questions for technical decision makers, data scientists, managers, developers, and VPs of engineering:</p>



(excluding c-suite) and professionals)	<ul style="list-style-type: none">• What customer experiences or other use cases do you expect to transform with AI?<ul style="list-style-type: none">○ How long do you have to wait for media to be created to finish building new products and user experiences?○ In what aspects of your supply chain are you experimenting with generative design technology?○ If you were able to use AI technology that is capable of generating media-rich content, how might that impact your digital assets strategy?• What have been some of your company's challenges in launching AI/ML-related products and business processes? Where do you see things getting blocked?<ul style="list-style-type: none">○ Do you have in-house development teams or outsource your development?○ How do you measure the developer productivity?○ Do you have any organizational goals around developer productivity or cost reduction?○ Have your software engineering teams begun experimenting with AI-generated code to accelerate rudimentary software engineering tasks?• Have you decided on which model you are using? Or do you plan to customize FMs for your particular use case or industry? <p>The following are examples of generative AI models:</p> <ul style="list-style-type: none">○ Amazon Titan (Titan Text and Titan Embeddings) from Amazon○ Jurassic-2 from AI21 Labs○ Claude from Anthropic○ Stable Diffusion from Stability AI○ Mistral from Mistral AI○ Llama from Meta <ul style="list-style-type: none">• Do you plan to customize FMs for your particular use case or industry?• Would you prefer to maintain control over your hosting instances?• Do you know how to evaluate the right model for the use case you are trying to solve with Generative AI?• What achievements or difficulties have you encountered when seeking assistance from an AWS partner to organize and develop your AI/ML capabilities?
Builders (Developer, data scientist, ML developer, ML engineer, BI engineer, director of development productivity, VP of engineering, software development managers)	<p>Builders are the ones implementing generative AI applications. They want to use foundation models to create generative AI applications aligned with business objectives. They are also searching for ways to enhance productivity of their teams to accelerate time to market for their applications.</p> <p>Discovery questions for data scientists and ML practitioners:</p> <ul style="list-style-type: none">• What ML platform are you currently using to evaluate and deploy custom ML models? What are some of the challenges that you face in the model lifecycle?

- Are you able to quickly and easily deploy foundation models today? What are the barriers?
- How are you partnering with your product and business stakeholders to provide them with generative AI capabilities quickly, securely, and cost effectively?

Look for the following:

- Data volume on or coming to AWS
- Create novel content such as images for ads, product descriptions, document summarizations, and contextual chatbots
- Experience using Amazon SageMaker
- Deployment of models on Amazon EC2 directly or from Hugging Face
- Desire to stay in AWS or customize a model for their domain (such as FinServ, or HCLS)

Discovery questions for developers:

- What is your team's current level of experience with coding and programming, and how much time do you typically spend on coding tasks?
- What type of application or project are you planning to use the code generator for?
- What specific coding tasks do you think a code generator would be most helpful for, and why?
- What are some examples of tasks or functions that you would like the code generator to automate for you?
 - How important is customizability to you, and how much flexibility do you need in terms of generating code that meets your specific needs?
 - What kind of support or training do you think you would need to use the code generator effectively, and what resources would be most helpful to you?
- Are there any particular features or functionalities that you would like the code generator to have, such as version control integration, testing or debugging support, or collaboration features?